

IML TensorFlow and Keras Workshop


Stefan Wunsch
stefan.wunsch@cern.ch

April 10, 2018

What is this workshop about?

- ▶ Modern description, implementation and application of neural networks
- ▶ Introduction to the currently favored packages:
 - ▶ **TensorFlow:** Low-level implementation of operations needed to implement neural networks in multi-threaded CPU and multi GPU environments
 - ▶ **Keras:** High-level convenience wrapper for backend libraries, e.g. TensorFlow, to implement neural network models



 [tensorflow / tensorflow](#)

 Watch ▾

7,714

★ Unstar

94,917

🔗 Fork

60,643

 [keras-team / keras](#)

 Watch ▾

1,621

★ Unstar

27,781

🔗 Fork

10,187

Outline

The workshop has these parts:

1. Very brief introduction to **neural networks**
2. Modern implementation of neural networks with **computational graphs** using **TensorFlow**
3. **Rapid development** of neural network applications using **Keras**

Assumptions of the tutorial:

- ▶ You are not a neural network expert, but you know roughly how to work.
- ▶ You don't know how TensorFlow and Keras works and how they play together.
- ▶ You want to know why TensorFlow and Keras are so popular and how you can use it!

Disclaimer:

- ▶ You won't learn how to use TensorFlow or Keras in one hour.
- ▶ **This tutorial tries to provide you with a good start and all information you need to become an expert!**

Set up your system

Clone the repository with the notebooks and slides:

```
git clone
```

```
https://github.com/stwunsch/iml_tensorflow_keras_workshop
```

Using SWAN (favored solution):

1. Log in on swan.cern.ch and select the software stack LCG 93
2. Open a terminal with New->Terminal and clone the repository as shown above
3. Browse to the notebooks as indicated on the following slides

Using your own laptop:


1. Clone the repository as shown above
2. Run the script `init_virtualenv.sh`
3. Source the virtual Python environment with `source py2_virtualenv/bin/activate`
4. Start a jupyter server with `jupyter notebook` and browse to the notebooks

Using lxplus:

1. Log in to lxplus with `ssh -Y your_username@lxplus.cern.ch`
2. Clone the repository as shown above
3. Source the software stack LCG 93 with `source /cvmfs/sft.cern.ch/lcg/views/LCG_93/x86_64-slc6-gcc62-opt/setup.sh`
4. Convert the notebooks to Python scripts with `jupyter nbconvert --to python input.ipynb output.py`

(Very) brief introduction to neural networks

A simple neural network




Neural Network: $f(x)$

- ▶ **Important:** A neural network is only a mathematical function. No magic involved!
- ▶ **Training:** Finding the best function for a given task, e.g., separation of signal and background.

Mathematical representation


- **Why do we need to know this?**

- TensorFlow implements these mathematical operations explicitly.
- Basic knowledge to understand Keras' high-level layers.



$$f_{\text{NN}} = \sigma(b_2 + W_2 \sigma(b_1 + W_1 x))$$

Mathematical representation (2)



$$f_{\text{NN}} = \sigma(b_2 + W_2 \sigma(b_1 + W_1 x))$$

$$\text{Input : } x = \begin{bmatrix} x_{1,1} \\ x_{2,1} \end{bmatrix}$$

$$\text{Weight : } W_1 = \begin{bmatrix} W_{1,1} & W_{1,2} \\ W_{2,1} & W_{2,2} \end{bmatrix}$$

$$\text{Bias : } b_1 = \begin{bmatrix} b_{1,1} \\ b_{2,1} \end{bmatrix}$$

Activation : $\sigma(x) = \tanh(x)$ (as example)

Activation is applied elementwise!

The “simple” neural network written as full equation:

$$f_{\text{NN}} = \sigma_2 \left([b_{1,1}^2] + [W_{1,1}^2 \quad W_{1,2}^2] \sigma_1 \left(\begin{bmatrix} b_{1,1}^1 \\ b_{2,1}^1 \end{bmatrix} + \begin{bmatrix} W_{1,1}^1 & W_{1,2}^1 \\ W_{2,1}^1 & W_{2,2}^1 \end{bmatrix} \begin{bmatrix} x_{1,1} \\ x_{2,1} \end{bmatrix} \right) \right)$$

Further reading: Deep Learning Textbook

Free textbook written by Ian Goodfellow, Yoshua Bengio and Aaron Courville:

<http://www.deeplearningbook.org/>

- ▶ Written by leading scientists in the field of machine learning
- ▶ **Everything you need to know** about modern machine learning and deep learning in particular.

▶ Part I: Applied Math and Machine Learning

Basics

- ▶ 2 Linear Algebra
- ▶ 3 Probability and Information Theory
- ▶ 4 Numerical Computation
- ▶ 5 Machine Learning Basics

▶ II: Modern Practical Deep Networks

- ▶ 6 Deep Feedforward Networks
- ▶ 7 Regularization for Deep Learning
- ▶ 8 Optimization for Training Deep Models
- ▶ 9 Convolutional Networks
- ▶ 10 Sequence Modeling: Recurrent and Recursive Nets
- ▶ 11 Practical Methodology
- ▶ 12 Applications

▶ III: Deep Learning Research

- ▶ 13 Linear Factor Models
- ▶ 14 Autoencoders
- ▶ 15 Representation Learning
- ▶ 16 Structured Probabilistic Models for Deep Learning
- ▶ 17 Monte Carlo Methods
- ▶ 18 Confronting the Partition Function
- ▶ 19 Approximate Inference
- ▶ 20 Deep Generative Models


Computational graphs with TensorFlow

What is TensorFlow?

TensorFlow is an open source software library for **numerical computation using data flow graphs**.


Nodes in the graph represent mathematical operations, while the graph edges represent the multidimensional data arrays (tensors) communicated between them.

- ▶ **In first place:** TensorFlow is not about neural networks.
- ▶ But it is a **perfect match** to implement neural networks efficiently!




$$f_{\text{NN}} = \sigma(b_2 + W_2 \sigma(b_1 + W_1 x))$$

Computational graphs



$$f_{NN} = \sigma(b_2 + W_2 \sigma(b_1 + W_1 x))$$



Example neural network



According computational graph


- ▶ TensorFlow implements all needed **mathematical operations for multi-threaded CPU and multi GPU** environments.
- ▶ Computation of neural networks using data flow graphs is a perfect match!

TensorFlow is an open source software library for numerical computation using data flow graphs. **Nodes** in the graph represent **mathematical operations**, while the **graph edges represent the multidimensional data arrays (tensors)** communicated between them.


Basic blocks to build graphs in TensorFlow

► Basic blocks:

- **Placeholders:** Used for injecting data into the graph, e.g., the inputs x of the neural network
- **Variables:** Free parameters of the graph, e.g., the weight matrices W of the neural network
- **Operations:** Functions that operate on data in the graph, e.g., the matrix multiplication of W_1 and x



$$f_{NN} = \sigma(b_2 + W_2 \sigma(b_1 + W_1 x))$$



Run the graph in a TensorFlow session

- ▶ A **graph** in TensorFlow can be run inside a **session**.
- ▶ Following example calculates $y = W \cdot x$ using TensorFlow:

Computational graph:

$$y = W \cdot x = \begin{pmatrix} 1 & 2 \end{pmatrix} \cdot \begin{pmatrix} 3 \\ 4 \end{pmatrix} = 11$$

TensorFlow code:

```
import tensorflow as tf
import numpy as np


# Build the graph y = W * x
x = tf.placeholder(tf.float32) # A placeholder
W = tf.get_variable("W", initializer=[[1.0, 2.0]]) # A variable
y = tf.matmul(W, x) # An operation

with tf.Session() as sess: # The session
    sess.run(tf.global_variables_initializer()) # Initialize variables
    result = sess.run(y, feed_dict={x: [[3.0], [4.0]]}) # Run graph
```

Example: XOR-solution with TensorFlow

Path to notebook: tensorflow/xor.ipynb

Scenario: Solving the separation of the blue crosses and red circles using a neural network implemented in TensorFlow




Content:

- ▶ Usage of placeholders, variables and operations to build a graph
- ▶ Run the graph in a session

Automatic differentiation

- ▶ XOR example covers only the inference (forward-pass) part of TensorFlow.
- ▶ Training includes optimization of weights using the back-propagation algorithm.
- ▶ Excessive use of gradients during training!

How can we compute the gradient of a graph?



Automatic differentiation (2)

- ▶ (Almost) each operation in TensorFlow is shipped with an inbuilt gradient.
- ▶ Computation of full gradient using the chain-rule of derivatives:

$$\frac{\partial z}{\partial x} = \frac{\partial z}{\partial y} \frac{\partial y}{\partial x}$$


- ▶ Explicit TensorFlow call: `tensorflow.gradients(z, x)`

Path to notebook:

`tensorflow/automatic_differentiation.ipynb`

Example: Full training tool-chain in TensorFlow

Path to notebook: tensorflow/gaussian.ipynb



Let's try to identify following steps:

1. Definition of neural network model
2. Implementation of loss function and optimizer algorithm
3. Training loop

Advanced: Efficient input pipelines in TensorFlow

- ▶ TensorFlow is designed to perform **highly-efficient computations** and ships **many useful features** ([Documentation](#)).
- ▶ Pick out one of the most-frequently needed: **Data-loading**
- ▶ Data-loading often bottleneck if not all data fits in memory (very common for image processing!)
- ▶ TensorFlow provides **input pipelines** directly **inbuilt in the graph**.
- ▶ **Full utilization of CPU/GPU** by loading data from disk in **queues** in memory concurrently

Path to notebook: `tensorflow/queues.ipynb`

And many more features ...

Further reading: Stanford course about TensorFlow

- ▶ Very well done and highly entertaining course!
- ▶ Lecturer working in the field (OpenAI, DeepMind, Google, ...)
- ▶ Small Keras part held by Francois Chollet (author of Keras!)

Link: <https://web.stanford.edu/class/cs20si/syllabus.html>

Rapid development of neural network applications using Keras

What is Keras?

- ▶ (Most) popular tool to train and apply neural networks
- ▶ **Python wrapper around multiple numerical computation libraries**, e.g., TensorFlow
- ▶ Hides most of the low-level operations that you don't want to care about.
- ▶ **Sacrificing little functionality** for much easier user interface
- ▶ **Backends:** TensorFlow, Theano, CNTK


Being able to go from idea to result with the least possible delay is key to doing good research.

theano




Why Keras and not one of the other wrappers?

- ▶ There are lot of alternatives: TFLearn, Lasagne, ...
- ▶ None of them are as **popular** as Keras!
- ▶ Keras is **tightly integrated into TensorFlow** and officially supported by Google.
- ▶ Looks like a **safe future for Keras!**



 keras-team / keras

Watch 1,621 Unstar 27,781 Fork 10,187

 kli-nlpr commented on Jan 16 Contributor +

Keras is gaining official Google support, and is moving into contrib, then core TF. If you want a high-level object-oriented TF API to use for the long term, Keras is the way to go.

<http://www.fast.ai/2017/01/03/keras/>

 1  7

- ▶ Read the full story here: [Link](http://www.fast.ai/2017/01/03/keras/)

Comparison of TensorFlow and Keras

Same model set up in TensorFlow and Keras:

TensorFlow:

```
def model(x):  
    with tf.variable_scope("model") as scope:  
        w1 = tf.get_variable('w1', shape=(2, 100), dtype=tf.float64,  
                              initializer=tf.random_normal_initializer(stddev=0.1))  
        b1 = tf.get_variable('b1', shape=(100), dtype=tf.float64,  
                              initializer=tf.constant_initializer(0.1))  
        w2 = tf.get_variable('w2', shape=(100, 1), dtype=tf.float64,  
                              initializer=tf.random_normal_initializer(stddev=0.1))  
        b2 = tf.get_variable('b2', shape=(1), dtype=tf.float64,  
                              initializer=tf.constant_initializer(0.1))  
  
        l1 = tf.nn.relu(tf.add(b1, tf.matmul(x, w1)))  
        logits = tf.add(b2, tf.matmul(l1, w2))  
    return logits, tf.sigmoid(logits)  
  
x = tf.placeholder(tf.float64, shape=[None, 2])  
logits, f = model(x)
```

Keras:

```
model = Sequential()  
model.add(Dense(100, activation="relu", input_dim=2))  
model.add(Dense(1, activation="sigmoid"))
```

Compare following notebooks for full code example:

Path to TensorFlow notebook: tensorflow/gaussian.ipynb

Path to Keras notebook: keras/gaussian.ipynb

Configure the Keras backend

Two ways to configure Keras backend (Theano, TensorFlow or CNTK):

1. Using **environment variables**
2. Using **Keras config file** in `$HOME/.keras/keras.json`

Example setup using environment variables:

Shell:

```
export KERAS_BACKEND=tensorflow
python your_script_using_keras.py
```

Inside a Python script:

```
from os import environ
environ['KERAS_BACKEND'] = 'tensorflow'
```

Example Keras config using TensorFlow as backend:

```
$ cat $HOME/.keras/keras.json
{
  "image_dim_ordering": "th",
  "epsilon": 1e-07,
  "floatx": "float32",
  "backend": "tensorflow"
}
```

Model definition with Keras

Path to notebook: keras/gaussian.ipynb

Model definition can be performed with two APIs:

Sequential model: Stacking layers sequentially

```
model = Sequential()  
model.add(Dense(100, activation="relu", input_dim=2))  
model.add(Dense(1, activation="sigmoid"))
```

Functional API: Multiple input/output models, ...

```
inputs = Input(shape=(2,))  
hidden_layer = Dense(100, activation="relu")(inputs)  
outputs = Dense(1, activation="sigmoid")(hidden_layer)  
model = Model(inputs=inputs, outputs=outputs)
```

```
model.summary()
```

Path to notebook: keras/gaussian.ipynb

Very useful convenience method:

```
model.summary()
```

```
=====
dense_1 (Dense)               (None, 100)                300
-----
dense_2 (Dense)               (None, 1)                   101
=====
Total params: 401
Trainable params: 401
Non-trainable params: 0
```

Easy to keep track of **model complexity**.

Setting optimizer, loss and validation metrics


Path to notebook: keras/gaussian.ipynb

Single line of code:

```
model.compile(  
    loss="binary_crossentropy", # Loss function  
    optimizer="adam",          # Optimizer algorithm  
    metrics=["accuracy"]       # Validation metric  
)
```

Available layers, losses, optimizers, ...

- ▶ There's **everything you can imagine**, and it's **well documented**.
- ▶ Have a look: www.keras.io



Training in Keras

Path to notebook: keras/gaussian.ipynb

Again, **single line of code:**

```
model.fit(data_train, labels_train,  
          validation_data=(data_val, labels_val),  
          batch_size=100,  
          epochs=100  
)
```

Save, load and apply the trained model

Save model:

- ▶ Models are **saved as HDF5 files**: `model.save("model.h5")`
 - ▶ Combines description of weights and architecture in a single file
- ▶ **Alternative**: Store weights and architecture separately
 - ▶ Store weights: `model.save_weights("model_weights.h5")`
 - ▶ Store architecture: `json_dict = model.to_json()`

Load model:

```
from keras.models import load_model  
model = load_model("model.h5")
```

Apply model:

```
predictions = model.predict(inputs)
```

Full example using the MNIST dataset

Path to notebook: `keras/mnist_train.ipynb`

► MNIST dataset?

- **Task:** Predict the number on an image of a handwritten digit
- **Official website:** Yann LeCun's website ([Link](#))
- Database of **70000 images of handwritten digits**
- 28x28 pixels in gray-scale as input, digit as label




► Data format:

- **Inputs:** 28x28 matrix with floats in $[0, 1]$
- **Target:** One-hot encoded digits, e.g., 2 \rightarrow $[0\ 0\ 1\ 0\ 0\ 0\ 0\ 0\ 0\ 0]$

Application on handwritten digits

Path to notebook: `keras/mnist_apply.ipynb`



If you are bored on your way home:

1. Open with GIMP `keras/your_own_digit.xcf`
2. Dig out your most beautiful handwriting
3. Save as PNG and run your model on it

Training with callbacks

Path to notebook: keras/mnist_train.ipynb

- ▶ **Callbacks** are executed before and/or after each training epoch.
- ▶ Numerous **predefined** callbacks are available, **custom** callbacks can be implemented.

Definition of model-checkpoint callback:

```
# Callback for model checkpoints  
checkpoint = ModelCheckpoint(  
    filepath="mnist_example.h5", # Output similar to model.save("mnist_example.h5")  
    save_best_only=True) # Save only model with smallest loss
```

Register callback:

```
model.fit(inputs, targets,  
    batch_size=100,  
    epochs=10,  
    callbacks=[checkpoint]) # Register callbacks
```

Training with callbacks (2)

Path to notebook: `keras/mnist_train.ipynb`

- ▶ Commonly used callbacks for improvement, debugging and validation of the training progress are implemented, e.g., **EarlyStopping**.
- ▶ Powerful tool: **TensorBoard** in combination with TensorFlow
- ▶ Custom callback: **LambdaCallback** or write callback class extending base class `keras.callbacks.Callback`

Callbacks

- Usage of callbacks
- Callback
- BaseLogger
- TerminateOnNaN
- ProgbarLogger
- History
- ModelCheckpoint
- EarlyStopping
- RemoteMonitor
- LearningRateScheduler
- TensorBoard
- ReduceLROnPlateau
- CSVLogger
- LambdaCallback
- Create a callback

Advanced: Customize Keras

Path to notebook:

keras/custom_loss_metric_callback.ipynb

- ▶ Keras is highly customizable!
- ▶ Easily define **own loss function, metrics and callbacks**

```
import keras.backend as K

def custom_loss(y_true, y_pred):
    return K.mean(K.square(y_pred - y_true), axis=-1)

def custom_metric(y_true, y_pred):
    return K.mean(K.square(y_pred - y_true), axis=-1)

model.compile(
    loss=custom_loss,
    metrics=[custom_metric],
    optimizer="adam")
```

Advanced: Training on “big data”

- ▶ The call `model.fit(inputs, targets, ...)` expects all inputs and targets to be already loaded in memory.
→ Physics applications have often data on Gigabyte to Terabyte scale!

These methods can be used to train on data that does not fit in memory.

- ▶ Training on **single batches**, performs a single gradient step:

```
model.train_on_batch(inputs, targets, ...)
```

- ▶ Training with data from a **Python generator**:

```
def generator_function():  
    while True:  
        inputs, labels = custom_load_next_batch()  
        yield inputs, labels
```

```
model.fit_generator(generator_function, ...)
```

Path to notebook: `keras/fit_generator.ipynb`

Note: Data preprocessing in Keras applications

- ▶ Some preprocessing methods are included in Keras, but mainly for text and image inputs.
- ▶ **Better option:** Using scikit-learn package ([Link to preprocessing module](#))

```
from sklearn.preprocessing import StandardScaler
preprocessing = StandardScaler()
preprocessing.fit(inputs)
preprocessed_inputs = preprocessing.transform(inputs)
```

Deep learning on the HIGGS dataset

One of the most often cited papers about deep learning in combination with a physics application:

Searching for Exotic Particles in High-Energy Physics with Deep Learning

Pierre Baldi, Peter Sadowski, Daniel Whiteson

- ▶ **Topic:** Application of deep neural networks for separation of signal and background in an exotic Higgs scenario
- ▶ **Results:** Deep learning neural networks are more powerful than “shallow” neural networks with only a single hidden layer.

Let's reproduce this with minimal effort using Keras!

Deep learning on the HIGGS dataset (2)

Path to notebook: keras/HIGGS_train.ipynb

Dataset:

- ▶ Number of events: 11M
- ▶ Number of features: 28

Shallow model:

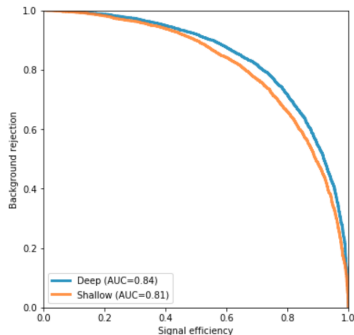

```
model_shallow = Sequential()  
model_shallow.add(Dense(1000, activation="tanh", input_dim=(28,)))  
model_shallow.add(Dense(1, activation="sigmoid"))
```

Deep model:

```
model_deep = Sequential()  
model_deep.add(Dense(300, activation="relu", input_dim=(28,)))  
model_deep.add(Dense(300, activation="relu"))  
model_deep.add(Dense(300, activation="relu"))  
model_deep.add(Dense(300, activation="relu"))  
model_deep.add(Dense(300, activation="relu"))  
model_deep.add(Dense(1, activation="sigmoid"))
```


Deep learning on the HIGGS dataset (3)

Path to notebook: `keras/HIGGS_test.ipynb`



- ▶ Shallow model matches performance in the paper, but deep model can be improved.
→ **Try to improve it!** But you'll need a decent GPU...
- ▶ Keras allows to **reproduce this result with a total of about 130 lines of code.**

Further reading: keras.io examples

- ▶ **Repository:** <http://www.github.com/keras-team/keras>
- ▶ **Folder:** `keras/examples/`

Recommendations:

- ▶ `addition_rnn.py`: Application of a RNN parsing strings such as "535+61" and returning the actual result 596, runs on a consumer CPU
- ▶ `neural_style_transfer.py`: Transfers the style of a reference image to an input image, needs a decent GPU

